

Deep Reinforcement Learning - Lab Assignment 1

Part #2 - Dynamic Programming (DP)

Autonomous Drone Rescue using Value Iteration

Team / Group Number (G): 178

We model an autonomous rescue drone as a finite **Markov Decision Process (MDP)** and solve it with **Value Iteration** to obtain the optimal value function $V^*(s)$ and the optimal policy $\pi^*(s)$.

Configuration derived from G = 178 (last digit 8):

- Last digit 8 is in 5-9 -> **6x6 grid, 3 rescue targets, 2 charging, 4 danger, 3 blocked**
- Last digit 8 is even -> **maximum battery = 10**
- Last digit 8 is in 5-9 -> **wind probability = 30%**
- Max step limit (6x6) -> **75 steps** (used for episode simulation)


```
In [1]: # Print execution timestamp and Virtual Machine ID at the top of the notebook
import datetime
import socket
import platform

print('Execution Timestamp :', datetime.datetime.now().strftime('%Y-%m-%d %H:%M:%S'))
print('Virtual Machine ID   :', socket.gethostname())
print('Platform             :', platform.platform())
print('Python Version       :', platform.python_version())
```

```
Execution Timestamp : 2026-06-08 04:04:27
Virtual Machine ID   : 33c34e5fc81f
Platform             : Linux-6.6.122+-x86_64-with-glibc2.35
Python Version       : 3.12.13
```

Virtual Lab Execution Screenshot (timestamped proof)

Below is a screenshot of this notebook (**Part #2 - DP**) executed inside the BITS virtual lab, showing the virtual-lab clock/timestamp (top-right) and the VM ID printed by the cell above.



1. Custom Drone Rescue Environment (1 Mark)

Grid configuration & symbol placement

Placement is generated deterministically from **G** (so it is reproducible and unique). The **start S is fixed at the top-left corner (0,0)**. Remaining special cells are drawn from a **G**-seeded shuffle of all other cells.

Symbol	Meaning
S	Start position (0,0)
F	Free / safe cell
D	Dangerous zone
R	Rescue target
C	Charging station
W	Wind zone (stochastic movement)
X	Blocked cell / obstacle

The assignment specifies counts for rescue/charging/danger/blocked cells but not for wind zones, so we add **2 wind zones** (documented assumption) to exercise the stochastic-transition logic.

The **+5** "reach charging station" reward is treated as a **one-time discovery bonus per station** (the battery still refills on every visit). Without this, repeatedly stepping onto a charger would form an infinite

positive-reward loop and the optimal policy would farm the charger forever instead of rescuing - so a per-station "discovered" flag is added to the state (documented assumption, analogous to a rescue target paying out once).

```
In [2]: import time
import random
import numpy as np
import matplotlib.pyplot as plt
from collections import defaultdict

G = 178 # group number
random.seed(G)
np.random.seed(G)

# ---- Parameters derived from the last digit of G (=8) ----
last_digit = G % 10
N = 6 if last_digit >= 5 else 5 # grid size (6x6 for digit 8)
MAX_BATTERY = 10 if last_digit % 2 == 0 else 15 # 10 (even last digit)
WIND_P = 0.30 if last_digit >= 5 else 0.20 # wind disturbance probability
MAX_STEPS = 75 if N == 6 else 50 # episode step cap (for simulation)

if last_digit >= 5:
    n_targets, n_charge, n_danger, n_block = 3, 2, 4, 3
else:
    n_targets, n_charge, n_danger, n_block = 2, 1, 3, 2
n_wind = 2 # documented assumption

# ---- Deterministic placement of special cells (seeded by G) ----
placer = random.Random(G) # dedicated RNG so placement is reproducible
free_cells = [(r, c) for r in range(N) for c in range(N) if (r, c) != (0, 0)]
placer.shuffle(free_cells)

rescue_targets = free_cells[0:n_targets]
charging_cells = free_cells[n_targets:n_targets + n_charge]
danger_cells = free_cells[n_targets + n_charge:n_targets + n_charge + n_danger]
blocked_cells = free_cells[n_targets + n_charge + n_danger:
                           n_targets + n_charge + n_danger + n_block]
wind_cells = free_cells[n_targets + n_charge + n_danger + n_block:
                       n_targets + n_charge + n_danger + n_block + n_wind]

# Sets for 0(1) membership tests
charging_set = set(charging_cells)
danger_set = set(danger_cells)
blocked_set = set(blocked_cells)
wind_set = set(wind_cells)
# Map each rescue-target cell to its index in the rescued-status tuple
target_index = {cell: i for i, cell in enumerate(rescue_targets)}
# Map each charging cell to its index in the charged-status tuple (one-time)
charge_index = {cell: i for i, cell in enumerate(charging_cells)}

START = (0, 0)
GAMMA = 0.95 # discount factor
THETA = 1e-3 # value-iteration stopping threshold
```

```

print(f'Grid size      : {N} x {N}')
print(f'Max battery    : {MAX_BATTERY}')
print(f'Wind probability : {WIND_P}')
print(f'Max steps/episode: {MAX_STEPS}')
print(f'Rescue targets   : {rescue_targets}')
print(f'Charging cells   : {charging_cells}')
print(f'Danger cells    : {danger_cells}')
print(f'Blocked cells    : {blocked_cells}')
print(f'Wind cells       : {wind_cells}')

```

```

Grid size      : 6 x 6
Max battery    : 10
Wind probability : 0.3
Max steps/episode: 75
Rescue targets : [(1, 1), (4, 4), (2, 0)]
Charging cells : [(4, 1), (0, 1)]
Danger cells   : [(3, 5), (0, 4), (3, 3), (0, 5)]
Blocked cells  : [(3, 2), (4, 2), (1, 5)]
Wind cells     : [(2, 2), (0, 2)]

```

```

In [3]: # ---- Action set and movement vectors ----
MOVES = ['Up', 'Down', 'Left', 'Right']           # the four movement actions
ACTIONS = MOVES + ['Hover']                       # full action space
DELTA = {'Up': (-1, 0), 'Down': (1, 0), 'Left': (0, -1), 'Right': (0, 1)}

def base_symbol(r, c):
    # Return the static map symbol for a cell (ignores the drone and rescue targets)
    if (r, c) == START: return 'S'
    if (r, c) in blocked_set: return 'X'
    if (r, c) in charging_set: return 'C'
    if (r, c) in danger_set: return 'D'
    if (r, c) in wind_set: return 'W'
    if (r, c) in target_index: return 'R'
    return 'F'

def move_cell(r, c, direction):
    # Apply a movement direction. The drone stays put if the target cell is
    # off-grid or blocked (the move still consumes battery, handled by the cost function)
    dr, dc = DELTA[direction]
    nr, nc = r + dr, c + dc
    if not (0 <= nr < N and 0 <= nc < N): # off-grid -> stay
        return (r, c)
    if (nr, nc) in blocked_set: # blocked -> stay
        return (r, c)
    return (nr, nc)

def valid_actions(state):
    # Return all valid actions from a state. Hover is always valid; a movement
    # is valid only if the adjacent cell is on the grid (moving into a blocked cell
    # is allowed but keeps the drone in place).
    r, c, b, resc, chg = state
    acts = ['Hover']
    for d in MOVES:
        dr, dc = DELTA[d]
        nr, nc = r + dr, c + dc

```

```

        if 0 <= nr < N and 0 <= nc < N:
            acts.append(d)
        return acts

# Print the symbolic grid layout
print('Initial grid layout:')
for r in range(N):
    print(' '.join(base_symbol(r, c) for c in range(N)))

```

Initial grid layout:

```

S C W F D D
F R F F F X
R F W F F F
F F X D F D
F C X F R F
F F F F F F

```

State representation & transition dynamics

State $s = (\text{row}, \text{col}, \text{battery}, \text{rescued}, \text{charged})$ where **rescued** is a tuple of booleans (one per rescue target) and **charged** is a tuple of booleans (one per charging station, marking whether its one-time +5 bonus has already been paid). This captures the drone position, remaining battery, which targets are still pending, and which chargers are already discovered - the information needed for the Markov property.

Rewards (mutually exclusive per step, matching the assignment table): rescue target +20, danger zone -10, charging station +5 (first visit only), regular movement -1, battery exhausted -20 (added when battery hits 0).

Battery: every action costs 1 unit; entering a charging station refills to full; hovering on a charging station adds +2 (capped at max); battery 0 ends the episode.

Blocked / off-grid moves: if a movement would leave the grid or enter a blocked cell the drone stays put and the step counts as a regular move (-1, battery -1); it does **not** re-collect the current cell's entry reward.

Wind: when the drone is on a wind cell and chooses a movement, with probability WIND_P the realised direction is replaced by a uniform random one of the four moves (so the intended direction keeps probability $1 - \text{WIND_P} + \text{WIND_P}/4$).

```

In [4]: def transitions(state, action):
        # Return the list of (probability, next_state, reward, done) tuples for
        # `action` in `state`. Assumes `state` is non-terminal.
        r, c, b, resc, chg = state
        out = defaultdict(float)                                # accumulate probability over ider

        def add(p, nstate, reward, done):
            out[(nstate, reward, done)] += p

        # ----- HOVER -----
        if action == 'Hover':

```

```

    if (r, c) in charging_set:          # hovering on a charger adds +2 ba
        nb = min(MAX_BATTERY, b + 2)
        add(1.0, (r, c, nb, resc, chg), 0.0, False)
    else:                                # hovering elsewhere costs 1 batte
        nb = b - 1
        done = (nb == 0)
        reward = -1 + (-20 if done else 0)
        add(1.0, (r, c, nb, resc, chg), reward, done)
    return [(p, ns, rew, dn) for (ns, rew, dn), p in out.items()]

# ----- MOVEMENT (with optional wind disturbance) -----
if (r, c) in wind_set:                  # build the (possibly randomised)
    dist = {}
    for d in MOVES:
        p = WIND_P / 4 + ((1 - WIND_P) if d == action else 0.0)
        dist[d] = dist.get(d, 0.0) + p
else:
    dist = {action: 1.0}

for d, p in dist.items():
    nr, nc = move_cell(r, c, d)         # resolve movement (handles walls/
    nb = b - 1                          # movement consumes 1 battery
    resc2 = list(resc)
    chg2 = list(chg)
    cell = (nr, nc)
    # Mutually-exclusive event reward for the entered cell
    if (nr, nc) == (r, c):
        # Move blocked / off-grid: the drone stays put. Per the assignme
        # consumes battery as a regular move - it must NOT re-trigger th
        # cell's entry reward (e.g. re-charging by bumping a wall while
        reward = -1
    elif cell in charging_set:
        nb = MAX_BATTERY                # charging always refills battery t
        ci = charge_index[cell]
        if not chg[ci]:                  # +5 incentive paid only the FIRST
            reward = 5                   # is reached (one-time, like a resc
            chg2[ci] = True              # mark this charger as already disc
        else:
            reward = 0                   # revisiting a known charger: refil
    elif cell in danger_set:
        reward = -10                     # danger penalty (does NOT terminat
    elif cell in target_index and not resc[target_index[cell]]:
        reward = 20                      # rescue reward
        resc2[target_index[cell]] = True # target now rescued (cell be
    else:
        reward = -1                      # regular movement cost
    done = False
    if nb == 0:                          # battery exhausted -> terminal wit
        reward += -20
        done = True
    resc2 = tuple(resc2)
    chg2 = tuple(chg2)
    if all(resc2):                       # all targets rescued -> terminal (
        done = True
    add(p, (nr, nc, nb, resc2, chg2), reward, done)

```

```
return [(p, ns, rew, dn) for (ns, rew, dn), p in out.items()]
```

```
In [5]: class DroneRescueEnv:
    # Gym-style environment exposing reset(), step(action) and render().
    def __init__(self):
        # Build the environment and initialise it to the starting state.
        self.reset()

    def reset(self):
        # Start at the top-left corner with a full battery, no targets rescued
        # no charging stations discovered yet.
        self.r, self.c = START
        self.b = MAX_BATTERY
        self.resc = tuple([False] * n_targets)
        self.chg = tuple([False] * n_charge)
        self.steps = 0
        return self.state()

    def state(self):
        # Return the current Markov state: (row, col, battery, rescued, charged)
        return (self.r, self.c, self.b, self.resc, self.chg)

    def step(self, action):
        # Sample one outcome from the transition distribution of the current state
        trans = transitions(self.state(), action)
        rnd = random.random()
        acc = 0.0
        chosen = trans[-1]
        for p, ns, rew, done in trans:
            acc += p
            if rnd <= acc:
                chosen = (p, ns, rew, done)
                break
        _, ns, reward, done = chosen
        self.r, self.c, self.b, self.resc, self.chg = ns
        self.steps += 1
        if self.steps >= MAX_STEPS:          # enforce the episode step cap
            done = True
        return ns, reward, done, {}

    def render(self):
        # Print the grid with the drone marked as 'A'; rescued targets shown in green
        print(f'Battery={self.b} Rescued={self.resc} Steps={self.steps}')
        for r in range(N):
            row = []
            for c in range(N):
                if (r, c) == (self.r, self.c):
                    row.append('A')                    # drone (agent)
                elif (r, c) in target_index and self.resc[target_index[(r, c)]]:
                    row.append('F')                    # already-rescued
                else:
                    row.append(base_symbol(r, c))
            print(' '.join(row))
        print()
```

```
# Quick sanity demo of the environment dynamics
env = DroneRescueEnv()
env.reset()
print('Initial environment state:')
env.render()
for a in ['Right', 'Down', 'Hover']:
    s, rew, done, _ = env.step(a)
    print(f'Action={a:5s} -> reward={rew}, done={done}')
env.render()
```

```
Initial environment state:
Battery=10  Rescued=(False, False, False)  Steps=0
A C W F D D
F R F F F X
R F W F F F
F F X D F D
F C X F R F
F F F F F F
```

```
Action=Right -> reward=5, done=False
Action=Down -> reward=20, done=False
Action=Hover -> reward=-1, done=False
Battery=8  Rescued=(True, False, False)  Steps=3
S C W F D D
F A F F F X
R F W F F F
F F X D F D
F C X F R F
F F F F F F
```

2. Dynamic Programming Solution - Value Iteration (2 Marks)

We enumerate all reachable states, then apply Value Iteration with discount `gamma = 0.95` and stopping threshold `theta = 1e-3`, reporting the number of iterations, runtime, and the final delta.

```
In [6]: def enumerate_states():
    # Enumerate every non-terminal state: a non-blocked position, battery in
    # any rescued-status combination that is not "all rescued" (terminal), a
    # charging-discovered combination (which chargers have already paid their
    states = []
    rescued_combos = []
    for mask in range(2 ** n_targets):    # all subsets of rescued targets
        combo = tuple(bool(mask & (1 << i)) for i in range(n_targets))
        rescued_combos.append(combo)
    charged_combos = []
    for mask in range(2 ** n_charge):    # all subsets of already-discovered
        combo = tuple(bool(mask & (1 << i)) for i in range(n_charge))
        charged_combos.append(combo)
    for r in range(N):
        for c in range(N):
            if (r, c) in blocked_set:    # drone can never occupy a blocked
```



```

        continue
    for b in range(1, MAX_BATTERY + 1): # battery 0 is terminal ->
        for combo in rescued_combos:
            if all(combo): # all-rescued is terminal -> exclu
                continue
            for ccombo in charged_combos:
                states.append((r, c, b, combo, ccombo))

    return states

all_states = enumerate_states()
print(f'Number of non-terminal states: {len(all_states)}')

```

Number of non-terminal states: 9240

```

In [7]: def value_iteration(gamma=GAMMA, theta=THETA):
    # Standard Value Iteration. V for terminal states stays 0 (default dict
    V = defaultdict(float)
    iterations = 0
    start_time = time.time()
    final_delta = 0.0
    while True:
        delta = 0.0
        for s in all_states:
            v_old = V[s]
            best = float('-inf')
            for a in valid_actions(s): # Bellman optimality backup
                q = 0.0
                for p, ns, reward, done in transitions(s, a):
                    q += p * (reward + (0.0 if done else gamma * V[ns]))
                best = max(best, q)
            V[s] = best
            delta = max(delta, abs(v_old - best))
        iterations += 1
        final_delta = delta
        # Per-iteration log so every Bellman sweep is visible in the PDF out
        print(f' Iteration {iterations:3d}: delta = {delta:.6f}')
        if delta < theta: # converged
            break
    runtime = time.time() - start_time
    return V, iterations, runtime, final_delta

V_star, n_iter, runtime, last_delta = value_iteration()
print(f'Converged in      : {n_iter} iterations')
print(f'Runtime          : {runtime:.3f} seconds')
print(f'Final delta/error  : {last_delta:.6f} (threshold = {THETA})')

```

```

Iteration 1: delta = 42.973747
Iteration 2: delta = 29.956112
Iteration 3: delta = 28.458307
Iteration 4: delta = 18.004875
Iteration 5: delta = 14.035551
Iteration 6: delta = 11.353676
Iteration 7: delta = 10.785992
Iteration 8: delta = 10.246692
Iteration 9: delta = 9.734358
Iteration 10: delta = 6.997832
Iteration 11: delta = 3.726613
Iteration 12: delta = 0.950000
Iteration 13: delta = 0.431780
Iteration 14: delta = 0.000341
Converged in      : 14 iterations
Runtime           : 5.859 seconds
Final delta/error : 0.000341 (threshold = 0.001)

```

```

In [8]: def greedy_policy(V, gamma=GAMMA):
        # Derive the optimal policy: the action maximising the one-step Bellman
        policy = {}
        for s in all_states:
            best_a, best_q = None, float('-inf')
            for a in valid_actions(s):
                q = 0.0
                for p, ns, reward, done in transitions(s, a):
                    q += p * (reward + (0.0 if done else gamma * V[ns]))
                if q > best_q:
                    best_q, best_a = q, a
            policy[s] = best_a
        return policy

pi_star = greedy_policy(V_star)
# Show the optimal action from the start state: full battery, nothing rescued
# no chargers discovered yet.
start_state = (0, 0, MAX_BATTERY, tuple([False] * n_targets), tuple([False]
print('Optimal action at start state', start_state, '->', pi_star[start_state])
print(f'V*(start) = {V_star[start_state]:.3f}')

```

```

Optimal action at start state (0, 0, 10, (False, False, False), (False, False, False)) -> Right
V*(start) = 51.117

```

3. Policy Visualisation (1 Mark)

The optimal policy is shown for **four representative state slices** so we can see how it adapts as battery and rescue progress change:

- **Full battery, none rescued** - the drone heads toward the nearest target sequence.
- **Full battery, only last target left** - the policy now routes to the remaining target.
- **Low battery, none rescued** - the policy prioritises charging stations before targets.
- **Low battery, only last target left** - same trade-off, but biased toward the final rescue.

Arrows show movement actions; a dot indicates Hover.

```
In [9]: # Arrow vectors for plotting movement directions (note: image y-axis points
        ARROW = {'Up': (0, 0.3), 'Down': (0, -0.3), 'Left': (-0.3, 0), 'Right': (0.3, 0)}


def draw_policy_panel(ax, batt, resc_mask, chg_mask, title):
    # Draw the optimal policy on one (battery, rescue, charge) slice.
    for r in range(N):
        for c in range(N):
            sym = base_symbol(r, c)
            # Colour-code special cells; rescued targets shown as free for c
            if sym == 'R' and resc_mask[target_index[(r, c)]]:
                cell_sym, color = 'F', 'white'
            else:
                cell_sym = sym
                color = {'X': '#444444', 'D': '#ff9999', 'C': '#99ccff',
                        'R': '#99ff99', 'W': '#ffe699', 'S': '#dddddd'}.get(sym, 'black')
            ax.add_patch(plt.Rectangle((c - 0.5, N - 1 - r - 0.5), 1, 1,
                                      facecolor=color, edgecolor='black'))
            ax.text(c, N - 1 - r + 0.32, cell_sym, ha='center', va='center',
                   fontsize=9, fontweight='bold')
            if sym == 'X':
                continue
            s = (r, c, batt, resc_mask, chg_mask)
            a = pi_star.get(s)
            if a == 'Hover':
                ax.plot(c, N - 1 - r, 'ko', markersize=5) # dot = hover
            elif a in ARROW:
                dx, dy = ARROW[a]
                ax.arrow(c, N - 1 - r, dx, dy, head_width=0.12,
                       head_length=0.12, fc='blue', ec='blue')
            ax.set_xlim(-0.5, N - 0.5); ax.set_ylim(-0.5, N - 0.5)
            ax.set_xticks(range(N)); ax.set_yticks(range(N))
            ax.set_title(title, fontsize=10)
            ax.set_aspect('equal')

# Four meaningful slices to show how the policy adapts to battery + rescue p
none_resc = tuple([False] * n_targets)
all_but_one = tuple([True] * (n_targets - 1) + [False]) # only last target
no_chg = tuple([False] * n_charge)
low_batt = max(1, MAX_BATTERY // 3) # ~1/3 battery slice

fig, axes = plt.subplots(2, 2, figsize=(12, 12))
draw_policy_panel(axes[0, 0], MAX_BATTERY, none_resc, no_chg,
                  f'Full battery ({MAX_BATTERY}), no targets rescued')
draw_policy_panel(axes[0, 1], MAX_BATTERY, all_but_one, no_chg,
                  f'Full battery ({MAX_BATTERY}), only last target left')
draw_policy_panel(axes[1, 0], low_batt, none_resc, no_chg,
                  f'Low battery ({low_batt}), no targets rescued')
draw_policy_panel(axes[1, 1], low_batt, all_but_one, no_chg,
                  f'Low battery ({low_batt}), only last target left')
fig.suptitle('Optimal Policy across battery / rescue slices', fontsize=12, f
plt.tight_layout()
plt.show()
```

Also keep the single-slice view for the heatmap cell to consume below

fixed_resc = none_resc
fixed_chg = no_chg
fixed_batt = MAX_BATTERY



4. State-Value Analysis (1 Mark)

We fix the rescue status (none rescued) and battery (full) and vary only the drone position, plotting a heatmap of $V^*(s)$.


```
In [10]: value_grid = np.full((N, N), np.nan) # NaN for blocked cells (shown b
for r in range(N):
    for c in range(N):
        if (r, c) in blocked_set:
            continue
        value_grid[r, c] = V_star[(r, c, fixed_batt, fixed_resc, fixed_chg)]
```

```

fig, ax = plt.subplots(figsize=(7, 6))
im = ax.imshow(value_grid, cmap='viridis')
plt.colorbar(im, ax=ax, label='V*(s)')
for r in range(N):
    for c in range(N):
        sym = base_symbol(r, c)
        txt = f'{sym}\n{value_grid[r, c]:.1f}' if not np.isnan(value_grid[r, c]) else ''
        ax.text(c, r, txt, ha='center', va='center', color='white', fontsize=10)
ax.set_title('State-Value Heatmap V*(s) (full battery, none rescued)')
ax.set_xticks(range(N))
ax.set_yticks(range(N))
plt.tight_layout()
plt.show()

print('Observation: cells closer to rescue targets (R) have higher V*, while')
print('danger zones (D) and far from any target have lower V*. Charging cell')
print('value high by protecting against battery exhaustion, and blocked cell')

```



Observation: cells closer to rescue targets (R) have higher V^* , while cells near danger zones (D) and far from any target have lower V^* . Charging cells (C) keep value high by protecting against battery exhaustion, and blocked cells are absent.

5. DP Scalability Discussion (1 Mark)

Curse of Dimensionality. The state space is the product of all state variables:

$\text{positions} \times \text{battery levels} \times 2^{(\#\text{targets})} \times 2^{(\#\text{chargers})}$. For our 6x6 grid this is roughly $33 \times 10 \times 2^3 \times 2^2$ states. Each variable we add multiplies the size.

How the state space grows:

- **10x10 grid:** positions jump from ~33 to ~100, roughly **3x** more states (and the reachable region grows), increasing both memory and per-sweep cost.
- **More rescue targets:** the rescued-status component is $2^{(\#\text{targets})}$, so each extra target **doubles** the state space (exponential growth).
- **Dynamic weather:** adding a wind/weather variable with w settings multiplies the state space by w and makes transitions time-dependent, breaking the stationary assumption DP relies on.

Is DP sufficient? DP gives an exact optimal policy and is ideal here because the MDP is small and fully known. But it requires (a) enumerating every state and (b) a known transition model. As dimensions grow, enumeration becomes infeasible (memory + time explode) and the model is often unknown in the real world.

How Deep RL helps. Methods like DQN / PPO replace the explicit value table with a neural-network function approximator that generalises across similar states, learn directly from sampled experience (no full model needed), and handle continuous or high-dimensional states - exactly the regime where tabular DP fails.

Real-world relation. A real rescue drone faces continuous positions, noisy sensors, unknown dynamics and changing weather; tabular DP cannot scale to this, so Deep RL (or model-based RL with learned models) is used in practice, with DP-style value iteration serving as the theoretical foundation and a baseline on small instances.