

Team 178 - MAB : Adaptive Treatment Recommendation (Multi-Armed Bandit)

Deep Reinforcement Learning - Lab Assignment 1

Part #1 - Multi-Armed Bandit (MAB)

Adaptive Treatment Recommendation System using Multi-Armed Bandit Learning

Team / Group Number (G): 178

Each medicine is modelled as an *arm* of a Multi-Armed Bandit. The system learns from patient outcomes over time and progressively identifies the optimal medicine.

Strategies implemented:

- **Greedy / Immediate Exploitation** (Task 2)
- **Epsilon-Greedy / Controlled Clinical Trial** (Task 3)
- **UCB1 / Confidence-Based** (Task 4)
- **Comparative Analysis** (Task 5)

Execution metadata (Virtual Lab requirement)

The cell below prints the **execution timestamp** and **Virtual Machine / Host ID**.

A timestamped screenshot from the virtual lab must accompany the final submission.

In [1]:

```
# Print execution timestamp and Virtual Machine ID at the top of the notebook
import datetime          # standard library for the current date/time
import socket            # used to read the machine/host name (acts as VM ID)
import platform          # extra environment details for traceability

# Current wall-clock time of execution (match this with the virtual-lab screenshot)
print('Execution Timestamp :', datetime.datetime.now().strftime('%Y-%m-%d %H:%M:%S'))
# Hostname acts as the Virtual Machine identifier inside the virtual lab
print('Virtual Machine ID  :', socket.gethostname())
# Platform info for full reproducibility context
print('Platform            :', platform.platform())
print('Python Version      :', platform.python_version())

Execution Timestamp : 2026-06-07 10:04:10
Virtual Machine ID  : Prithvis-MacBook-Air.local
Platform           : macOS-26.3-arm64-arm-64bit-Mach-O
Python Version     : 3.13.8
```

Virtual Lab Execution Screenshot (timestamped proof)

Below is a screenshot of this notebook (**Part #1 - MAB**) executed inside the BITS virtual lab, showing the virtual-lab clock/timestamp (top-right) and the VM ID printed by the cell above.

![Virtual Lab Execution Screenshot](attachment:virtual_lab_MAB.png)

Task 1: Dataset Design (1 Mark)

The synthetic patient-treatment environment is derived from the group number $G = 178$ so the dataset is unique and reproducible.

Derivation for $G = 178$

- Medicines: $K = (G \bmod 3) + 5 = 1 + 5 = 6$
- Hidden success probability: $P_i = 0.4 + ((G + i) \bmod 6) * 0.07$
- Severity: $Severity = (patient_id \bmod 5) + 1$ (1 = mild ... 5 = critical)
- Utility (reward): $utility = clinical_outcome * (1 - Severity/10)$

In [2]:

```
# ---- Core imports ----
import random                # python RNG (seeded for reproducibility)
import numpy as np           # vectorised maths and RNG
import pandas as pd          # tabular dataset handling
import matplotlib.pyplot as plt # plotting cumulative-reward curves

# ---- Reproducibility: seed both RNGs with the group number ----
G = 178                      # our group number
random.seed(G)               # seed python's random module
np.random.seed(G)            # seed numpy's global RNG

# ---- Environment parameters derived from G ----
K = (G % 3) + 5               # number of medicines (arms)
# Hidden success probability for each medicine i in {0,...,K-1}
true_probs = np.array([0.4 + ((G + i) % 6) * 0.07 for i in range(K)])

N_PATIENTS = 1000            # total patients (iterations) to simulate

print('Group number G      : ', G)
print('Total medicines K   : ', K)
print('Hidden success probabilities (P_i):')
for i, p in enumerate(true_probs):
    print(f'  Medicine {i}: P = {p:.2f}')
print(f'\nTrue best medicine : {int(np.argmax(true_probs))} (P = {true_probs.max():.2f})')

Group number G      : 178
Total medicines K   : 6
Hidden success probabilities (P_i):
  Medicine 0: P = 0.68
  Medicine 1: P = 0.75
  Medicine 2: P = 0.40
  Medicine 3: P = 0.47
  Medicine 4: P = 0.54
  Medicine 5: P = 0.61

True best medicine : 1 (P = 0.75)
```

In [3]:

```
# ---- Build the static part of the dataset (patient_id + severity_score) ----
# assigned_medicine / clinical_outcome / utility_score are populated *dynamically*
# by each algorithm at run time, so the base table only holds the fixed fields.
patient_ids = np.arange(N_PATIENTS)          # 0 .. 999
severity_scores = (patient_ids % 5) + 1       # severity in range 1..5

dataset = pd.DataFrame({
    'patient_id': patient_ids,                # patient index
    'severity_score': severity_scores,         # disease severity (1-5)
    'assigned_medicine': np.nan,               # filled during an algorithm run
    'clinical_outcome': np.nan,               # filled during an algorithm run
    'utility_score': np.nan                   # filled during an algorithm run
})
```

```
print('First 10 dataset rows:')
print(dataset.head(10).to_string(index=False))
```

First 10 dataset rows:

patient_id	severity_score	assigned_medicine	clinical_outcome	utility_score
0	1	NaN	NaN	NaN
1	2	NaN	NaN	NaN
2	3	NaN	NaN	NaN
3	4	NaN	NaN	NaN
4	5	NaN	NaN	NaN
5	1	NaN	NaN	NaN
6	2	NaN	NaN	NaN
7	3	NaN	NaN	NaN
8	4	NaN	NaN	NaN
9	5	NaN	NaN	NaN

In [4]:

```
def administer_treatment(medicine, severity):
    # Simulate giving `medicine` to a patient of the given `severity`.
    # Returns:
    #   clinical_outcome (int) : 1 if patient recovers else 0 (Bernoulli with P_i)
    #   utility_score (float) : reward = clinical_outcome * (1 - severity/10)
    # clinical_outcome is used to UPDATE the bandit estimates;
    # utility_score is used to ACCUMULATE the cumulative reward.
    clinical_outcome = 1 if np.random.random() < true_probs[medicine] else 0
    utility_score = clinical_outcome * (1 - severity / 10)
    return clinical_outcome, utility_score
```

Task 2: Immediate Exploitation Strategy (1 Mark)

Policy: *"Once a treatment appears best, keep prescribing only that treatment."*

- **Warm-up:** test each medicine exactly **10 times** ($10 * K$ patients).
- **Exploit:** afterwards always pick the medicine with the highest estimated recovery rate ($\text{argmax } Q$).

In [5]:

```
def run_greedy(initial_pulls=10):
    # Greedy / immediate-exploitation strategy over N_PATIENTS patients.
    # Each medicine is tried `initial_pulls` times (warm-up); then the best-so-far
    # medicine is always selected.
    np.random.seed(G) # re-seed: every strategy faces same outcomes
    Q = np.zeros(K) # estimated recovery rate per medicine
    N = np.zeros(K) # number of times each medicine was used
    df = dataset.copy() # private copy of the dataset to populate
    cumulative = np.zeros(N_PATIENTS) # cumulative reward after each patient
    best_hist = np.zeros(N_PATIENTS, int) # running best-estimated arm after each patient
    total = 0.0

    for t in range(N_PATIENTS):
        severity = int(df.at[t, 'severity_score'])
        if t < initial_pulls * K: # warm-up: round-robin over all medicines
            arm = t % K
        else: # exploit: pick best estimated medicine
            arm = int(np.argmax(Q))
        outcome, utility = administer_treatment(arm, severity)
        N[arm] += 1 # update usage count
```

```

    Q[arm] += (outcome - Q[arm]) / N[arm]    # incremental sample-average update
    total += utility                        # accumulate utility reward
    df.at[t, 'assigned_medicine'] = arm
    df.at[t, 'clinical_outcome'] = outcome
    df.at[t, 'utility_score'] = utility
    cumulative[t] = total
    best_hist[t] = int(np.argmax(Q))    # which medicine currently looks best
return {'name': 'Greedy', 'df': df, 'Q': Q, 'N': N,
        'cumulative': cumulative, 'best_hist': best_hist, 'total': total}

```

```

greedy_res = run_greedy(initial_pulls=10)
print('Greedy estimated recovery rates Q:', np.round(greedy_res['Q'], 3))
print('Greedy pulls per medicine      : ', greedy_res['N'].astype(int))
print('Greedy chosen best medicine    : ', int(np.argmax(greedy_res['Q'])))
print(f"Greedy cumulative reward (1000) : {greedy_res['total']:.2f}")
print('\nFirst 10 populated rows:')
print(greedy_res['df'].head(10).to_string(index=False))

```

```

Greedy estimated recovery rates Q: [0.7  0.755 0.2  0.4  0.4  0.667]
Greedy pulls per medicine      : [ 10 948  10  10  10  12]
Greedy chosen best medicine    : 1
Greedy cumulative reward (1000) : 520.10

```

First 10 populated rows:

patient_id	severity_score	assigned_medicine	clinical_outcome	utility_score
0	1	0.0	1.0	0.9
1	2	1.0	1.0	0.8
2	3	2.0	0.0	0.0
3	4	3.0	0.0	0.0
4	5	4.0	0.0	0.0
5	1	5.0	0.0	0.0
6	2	0.0	1.0	0.8
7	3	1.0	0.0	0.0
8	4	2.0	0.0	0.0
9	5	3.0	1.0	0.5

Task 3: Controlled Clinical Trial - Epsilon-Greedy (1.5 Marks)

Policy: mostly give the current best treatment, but with probability **epsilon** explore a random treatment to discover hidden opportunities.

Main run uses **epsilon** = **0.10**; we also analyse **0.01** and **0.50**.

In [6]:

```

def run_epsilon_greedy(epsilon):
    # Epsilon-greedy strategy: explore a random medicine with probability `epsilon`,
    # otherwise exploit the medicine with the highest estimated recovery rate.
    np.random.seed(G)                # same outcome stream for a fair comparison
    Q = np.zeros(K)
    N = np.zeros(K)
    df = dataset.copy()
    cumulative = np.zeros(N_PATIENTS)
    best_hist = np.zeros(N_PATIENTS, int)
    total = 0.0

    for t in range(N_PATIENTS):
        severity = int(df.at[t, 'severity_score'])
        if np.random.random() < epsilon:    # explore: random medicine

```

```

        arm = np.random.randint(K)
    else:
        # exploit: best estimated medicine
        arm = int(np.argmax(Q))
    outcome, utility = administer_treatment(arm, severity)
    N[arm] += 1
    Q[arm] += (outcome - Q[arm]) / N[arm]
    total += utility
    df.at[t, 'assigned_medicine'] = arm
    df.at[t, 'clinical_outcome'] = outcome
    df.at[t, 'utility_score'] = utility
    cumulative[t] = total
    best_hist[t] = int(np.argmax(Q))
return {'name': f'Eps-Greedy ({epsilon:.0%})', 'df': df, 'Q': Q, 'N': N,
        'cumulative': cumulative, 'best_hist': best_hist, 'total': total}

```

```

eps10_res = run_epsilon_greedy(0.10)      # main run: 10% exploration
print('Epsilon = 10%')
print(' Estimated Q      : ', np.round(eps10_res['Q'], 3))
print(' Pulls per medicine: ', eps10_res['N'].astype(int))
print(f" Cumulative reward : {eps10_res['total']:.2f}")
print(' First 10 populated rows:')
print(eps10_res['df'].head(10).to_string(index=False))

```

Epsilon = 10%

Estimated Q : [0.476 0.734 0.296 0.412 0.636 0.462]

Pulls per medicine: [21 887 27 17 22 26]

Cumulative reward : 494.10

First 10 populated rows:

patient_id	severity_score	assigned_medicine	clinical_outcome	utility_score
0	1	0.0	1.0	0.9
1	2	0.0	0.0	0.0
2	3	0.0	0.0	0.0
3	4	0.0	0.0	0.0
4	5	0.0	1.0	0.5
5	1	1.0	1.0	0.9
6	2	1.0	1.0	0.8
7	3	1.0	1.0	0.7
8	4	5.0	0.0	0.0
9	5	1.0	0.0	0.0

In [7]:

---- Sensitivity analysis: epsilon = 1%, 10%, 50% ----

eps01_res = run_epsilon_greedy(0.01) # very little exploration

eps50_res = run_epsilon_greedy(0.50) # heavy exploration

print('Effect of exploration rate on final cumulative reward:')

for res in (eps01_res, eps10_res, eps50_res):

best = int(np.argmax(res['Q']))

print(f" {res['name']:>18s} -> reward = {res['total']:.2f} | best medicine = {best}")

print('\nObservation:')

print(' * 1% -> too little exploration; risks locking onto a sub-optimal medicine.')

print(' * 10% -> balanced; quickly finds the best medicine and mostly exploits it.')

print(' * 50% -> too much exploration; wastes ~half the patients on inferior medicines.')

Effect of exploration rate on final cumulative reward:

Eps-Greedy (1%) -> reward = 478.90 | best medicine = 0

Eps-Greedy (10%) -> reward = 494.10 | best medicine = 1

Eps-Greedy (50%) -> reward = 456.00 | best medicine = 1

Observation:

- * 1% -> too little exploration; risks locking onto a sub-optimal medicine.
- * 10% -> balanced; quickly finds the best medicine and mostly exploits it.
- * 50% -> too much exploration; wastes ~half the patients on inferior medicines.

Task 4: Confidence-Based Strategy - UCB1 (1 Mark)

Select the arm maximising the UCB1 index:

$$a_t = \arg\max_i \left(Q_i + \sqrt{\frac{2 \ln t}{N_i}} \right)$$

Each medicine is tried once first so every $N_i > 0$ before the bonus is used.

In [8]:

```
def run_ucb1():
    # UCB1 confidence-based strategy. Medicines with fewer observations receive a
    # larger exploration bonus that shrinks as more evidence is collected.
    np.random.seed(G)
    Q = np.zeros(K)
    N = np.zeros(K)
    df = dataset.copy()
    cumulative = np.zeros(N_PATIENTS)
    best_hist = np.zeros(N_PATIENTS, int)
    total = 0.0

    for t in range(N_PATIENTS):
        severity = int(df.at[t, 'severity_score'])
        if t < K:
            # initialise: try each medicine once
            arm = t
        else:
            # pick the arm with the largest UCB1 index
            ucb_values = Q + np.sqrt(2 * np.log(t + 1) / N)
            arm = int(np.argmax(ucb_values))
        outcome, utility = administer_treatment(arm, severity)
        N[arm] += 1
        Q[arm] += (outcome - Q[arm]) / N[arm]
        total += utility
        df.at[t, 'assigned_medicine'] = arm
        df.at[t, 'clinical_outcome'] = outcome
        df.at[t, 'utility_score'] = utility
        cumulative[t] = total
        best_hist[t] = int(np.argmax(Q))
    return {'name': 'UCB1', 'df': df, 'Q': Q, 'N': N,
            'cumulative': cumulative, 'best_hist': best_hist, 'total': total}

ucb_res = run_ucb1()
print('UCB1 estimated recovery rates Q:', np.round(ucb_res['Q'], 3))
print('UCB1 pulls per medicine      :', ucb_res['N'].astype(int))
print('UCB1 chosen best medicine    :', int(np.argmax(ucb_res['Q'])))
print(f'UCB1 cumulative reward (1000) : {ucb_res['total']:.2f}')
print('\nFirst 10 populated rows:')
print(ucb_res['df'].head(10).to_string(index=False))
```

```
UCB1 estimated recovery rates Q: [0.673 0.767 0.478 0.426 0.46  0.584]
UCB1 pulls per medicine      : [196 507  67  54  63 113]
UCB1 chosen best medicine    : 1
UCB1 cumulative reward (1000) : 469.90
```

First 10 populated rows:

patient_id	severity_score	assigned_medicine	clinical_outcome	utility_score
0	1	0.0	1.0	0.9
1	2	1.0	1.0	0.8
2	3	2.0	0.0	0.0
3	4	3.0	0.0	0.0
4	5	4.0	0.0	0.0
5	1	5.0	0.0	0.0
6	2	0.0	1.0	0.8
7	3	1.0	0.0	0.0
8	4	0.0	0.0	0.0
9	5	2.0	1.0	0.5

Task 5: Comparative Analysis (0.5 Mark)


Plot **Cumulative Reward vs. Number of Patients** for every strategy.

In [9]:

```
# Compare every strategy on one cumulative-reward plot
strategies = [greedy_res, eps01_res, eps10_res, eps50_res, ucb_res]

plt.figure(figsize=(10, 6))
x = np.arange(1, N_PATIENTS + 1)          # patient index on the x-axis
for res in strategies:                     # one curve per strategy
    plt.plot(x, res['cumulative'], label=f"{res['name']} ({res['total']:.1f})")
plt.xlabel('Number of Patients')
plt.ylabel('Cumulative Reward (utility)')
plt.title('Cumulative Reward vs. Number of Patients (Group 178)')
plt.legend()
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()

print('Final cumulative reward ranking:')
for res in sorted(strategies, key=lambda r: r['total'], reverse=True):
    print(f" {res['name']:>18s} : {res['total']:7.2f}")
```



Final cumulative reward ranking:

```

Greedy : 520.10
Eps-Greedy (10%) : 494.10
Eps-Greedy (1%) : 478.90
UCB1 : 469.90
Eps-Greedy (50%) : 456.00

```

Quantitative answers computed from the run

The cell below derives the answers to Questions 1-3 **directly from this run's data** so the stated conclusions always match the numbers above.

- **Highest cumulative reward** = strategy with the largest final **total**.
- **Convergence - two metrics:**
 - ***Loose:*** earliest patient after which the running best-estimated arm stays constant.

This can trigger early simply because a high-exploration strategy samples every arm quickly, so it is reported for transparency but **not** used for the final answer.

- ***Steady-state lock-in (clinically meaningful):*** earliest patient **T** such that for the next 100 patients, the strategy **actually chose** the true best medicine at least 90% of the time. This captures both ***identification*** and ***consistent exploitation***; strategies that keep exploring forever (e.g. Eps-Greedy 50%) never reach this threshold and are reported as **never**.

- **Most stable** = smallest standard deviation of the per-patient reward over the last 200 patients (fewest fluctuations).

In [10]:

```

true_best = int(np.argmax(true_probs))      # the genuinely optimal medicine (=1)

def loose_convergence(res):
    # LOOSE metric: earliest patient index after which the running best-estimated arm
    # stays constant for the rest of the run. Reported for transparency only -- a
    # heavy explorer like Eps-Greedy(50%) trivially satisfies this within ~K patients
    # because it samples every arm early, even though it keeps randomly switching
    # ACTIONS forever.

```

```

bh = res['best_hist']
final = bh[-1]
conv = 0
for t in range(N_PATIENTS - 1, -1, -1):
    if bh[t] == final:
        conv = t
    else:
        break
return conv

def steady_state_lock(res, window=100, threshold=0.9):
    # CLINICALLY MEANINGFUL convergence metric:
    # Earliest patient index T such that, for the window [T, T+window-1], the strategy
    # actually CHOSE the true best medicine at least `threshold` of the time. This
    # captures both (a) the running estimate has identified the true best AND (b) the
    # strategy is consistently exploiting it. Pure-random / high-eps strategies that
    # keep exploring will never reach 90% true-best usage and are reported as "never".
    actions = res['df']['assigned_medicine'].values.astype(int)
    n = len(actions)
    for t in range(n - window + 1):
        win = actions[t:t + window]
        if (win == true_best).mean() >= threshold:
            return t
    return -1 # never reached steady state on the true best

def stability(res):
    # Std-dev of per-patient reward over the last 200 patients (lower = more stable).
    increments = np.diff(np.concatenate([[0.0], res['cumulative']]))
    return float(np.std(increments[-200:]))

print('Strategy          | final reward | best medicine | loose-conv@ | steady-lock@(\u226590%
true-best in 100-window) | stability(std)')
for res in strategies:
    sl = steady_state_lock(res)
    sl_str = f'{sl:^46d}' if sl >= 0 else f'{'never':^46s}'
    print(f" {res['name']:>16s} | {res['total']:11.2f} | "
          f"{int(np.argmax(res['Q'])):^13d} | {loose_convergence(res):^11d} | "
          f"{sl_str} | {stability(res):.4f}")

highest = max(strategies, key=lambda r: r['total'])
# For "fastest convergence" use the steady-state metric: skip strategies that never lock in.
locked = [r for r in strategies if steady_state_lock(r) >= 0]
fastest = min(locked, key=steady_state_lock) if locked else min(strategies, key=loose_convergence)
most_stable = min(strategies, key=stability)
print('\nQ1 Highest cumulative reward:', highest['name'])
print('Q2 Fastest convergence     :', fastest['name'],
      '(first patient at which the next 100-window chooses TRUE best >=90%% =',
      steady_state_lock(fastest), ')')
print('Q3 Most stable performance :', most_stable['name'])
print(' True best medicine         :', true_best,
      '(P =', round(float(true_probs[true_best]), 2), ')')

Strategy          | final reward | best medicine | loose-conv@ | steady-lock@(\u226590% true-best in
100-window) | stability(std)
Greedy |      520.10 |      1      |      120      |
| 0.3212

```

Eps-Greedy (1%)		478.90		0		685		never
0.3518								
Eps-Greedy (10%)		494.10		1		644		4
0.3486								
Eps-Greedy (50%)		456.00		1		13		never
0.3465								
UCB1		469.90		1		18		never
0.3344								

Q1 Highest cumulative reward : Greedy

Q2 Fastest convergence : Eps-Greedy (10%) (first patient at which the next 100-window chooses TRUE best $\geq 90\%$ = 4)

Q3 Most stable performance : Greedy

True best medicine : 1 (P = 0.75)

Analysis Questions & Comparative Summary

Questions 1-3 are answered by the computed cell above (so they always match this run). The discussion below interprets those results for **G = 178**.

1. Highest cumulative reward at 1000 patients?

See **Q1** above - for this run it is **Greedy (~520)**. With a clean round-robin warm-up (10 pulls each) the true best medicine (medicine 1, P = 0.75) stands out clearly, so Greedy locks onto it and then spends *every* remaining patient exploiting it - wasting nothing on exploration. Epsilon-Greedy (10%) is a close second; Epsilon-Greedy (50%) is lowest because it keeps prescribing random (often inferior) medicines.

2. Fastest convergence?

We deliberately use the **steady-state lock-in** metric (see explanation above) rather than the loose "running best stays fixed" criterion, because the latter is misleading: a heavy explorer like Epsilon-Greedy (50%) trivially satisfies it within $\sim K$ patients just by sampling every arm, even though it then keeps randomly switching its *actions* forever (only $\sim 50\%$ of its decisions match the true best). The steady-state metric requires the strategy to *both* identify the true best *and* consistently prescribe it ($\geq 90\%$ of the next 100 patients) - exactly what "converged" should mean clinically. By this metric the fastest is **Epsilon-Greedy (10%)** (locks in around patient ~ 4), with **Greedy** following once its 60-patient round-robin warm-up completes ($\sim$ patient 51). Epsilon-Greedy (50%) and UCB1 are reported as **never** here: $\epsilon=0.50$ keeps exploring on half the patients so it caps at $\sim 58\%$ true-best usage, and UCB1 keeps "owing" exploration to under-pulled arms throughout the 1000-patient horizon (it does prefer medicine 1 but stays below the strict 90% bar in any 100-window).

3. Most stable performance?

See **Q3** above - **Greedy** is the most stable: once it commits to a single medicine the only remaining variation is that medicine's intrinsic Bernoulli/severity noise, giving the smoothest curve. Epsilon-Greedy (50%) fluctuates the most due to constant random switching.

4. Safest approach for real-world hospital deployment?

UCB1. It is deterministic (no random prescriptions on real patients), gives principled extra chances to under-tested treatments, and offers a confidence-based justification for every choice - all valuable when patient safety matters. Pure Greedy "won" on reward here only because of a lucky, clean warm-up; on a different seed a poor warm-up can permanently lock it onto a sub-optimal medicine, which is unacceptable clinically. Epsilon-Greedy is attractive operationally (it *did* lock in fastest by our strict metric) but it never fully stops issuing random prescriptions, which is also clinically unattractive. UCB1's shrinking confidence bonus gives the best of both: principled exploration that tapers without ever being purely random.

Comparative summary.

On this dataset Greedy achieves the highest cumulative reward and the most stable curve because the warm-up cleanly identifies medicine 1 and it then exploits without waste. Epsilon-Greedy trades a little reward for robustness: 10% performs well and locks onto the true best fastest by the strict metric, 1% risks getting stuck on a sub-optimal arm (here it actually does, settling on medicine 0), and 50% wastes too many patients on inferior medicines. UCB1 evaluates every treatment fairly with an explicit confidence bonus; although its raw reward is slightly lower over 1000 patients, its determinism and built-in confidence guarantees make it the recommended, safest choice for a real clinical setting where prescribing randomly is unacceptable.